

Classification Report: Predicting Crime Severity

```
In [413]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report, r2_score
from statsmodels.discrete.discrete_model import Logit
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
```

```
In [414]: crime = pd.read_csv("Crime_Data_from_2020_to_Present.csv")
print(crime.shape)

(1112545, 28)
```

```
In [415]: crime.head()
```

Out [415]:

	DR_NO	Date Rptd	DATE OCC	TIME OCC	AREA	AREA NAME	Rpt Dist No	Part 1-2	Crm Cd	Crm Cd D
0	10304468	01/08/2020 12:00:00 AM	01/08/2020 12:00:00 AM	2230	3	Southwest	377	2	624	BATTEF SIMI ASSAULT
1	190101086	01/02/2020 12:00:00 AM	01/01/2020 12:00:00 AM	330	1	Central	163	2	624	BATTEF SIMI ASSAULT
2	200110444	04/14/2020 12:00:00 AM	02/13/2020 12:00:00 AM	1200	1	Central	155	2	845	SEX OFFENSE REGISTRATION OUT COMPLIANCE
3	191501505	01/01/2020 12:00:00 AM	01/01/2020 12:00:00 AM	1730	15	N Hollywood	1543	2	745	VANDALISM MISDEAMEANOR (\$399 UNDERTH
4	191921269	01/01/2020 12:00:00 AM	01/01/2020 12:00:00 AM	415	19	Mission	1998	2	740	VANDALISM FELONY (\$400 OVER, 1 CHURCH V

5 rows × 28 columns

```
In [416]: crime.columns
```


```
Out[416]: Index(['DR_NO', 'Date Rptd', 'DATE OCC', 'TIME OCC', 'AREA', 'AREA N  
AME',  
                'Rpt Dist No', 'Part 1-2', 'Crm Cd', 'Crm Cd Desc', 'Mocode  
s',  
                'Vict Age', 'Vict Sex', 'Vict Descent', 'Premis Cd', 'Premis  
Desc',  
                'Weapon Used Cd', 'Weapon Desc', 'Status', 'Status Desc', 'Cr  
m Cd 1',  
                'Crm Cd 2', 'Crm Cd 3', 'Crm Cd 4', 'LOCATION', 'Cross Stree  
t', 'LAT',  
                'LON'],  
              dtype='object')
```

```
In [417]: # Data cleaning
```

```
crime = crime.drop_duplicates()  
print(crime.shape)  
  
crime['AREA'] = pd.to_numeric(crime['AREA'], errors='coerce')  
print(crime['AREA'].dtype)
```

```
(881775, 28)  
int64
```


```
In [418]: # Example: Crime Distribution by LA Area
plt.figure(figsize=(14, 10))
ax=sns.countplot(x="AREA NAME", data=crime)
plt.title("Crime Distribution by LA Area")
plt.xticks(rotation=90)
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2.,
                                     ha='center', va='center', xytext=(0, 6), textcoords='c
plt.show()
```



```
In [422]: print(crime['Crm Cd'].describe())

plt.subplot(1, 2, 2)
plt.hist(crime['Crm Cd'], color='blue')
plt.title("Density of crime severity for 2020 - Present")
plt.show()
```

```
count      881775.000000
mean         500.990648
std          207.841176
min          110.000000
25%          331.000000
50%          442.000000
75%          626.000000
max          956.000000
Name: Crm Cd, dtype: float64
```




In [423]: `crime.columns`

Out[423]: Index(['DR_NO', 'Date Rptd', 'DATE OCC', 'TIME OCC', 'AREA', 'AREA NAME',
 'Rpt Dist No', 'Part 1-2', 'Crm Cd', 'Crm Cd Desc', 'Mocode s',
 'Vict Age', 'Vict Sex', 'Vict Descent', 'Premis Cd', 'Premis Desc',
 'Weapon Used Cd', 'Weapon Desc', 'Status', 'Status Desc', 'Crm Cd 1',
 'Crm Cd 2', 'Crm Cd 3', 'Crm Cd 4', 'LOCATION', 'Cross Street', 'LAT',
 'LON'],
 dtype='object')

In [424]: `# Renaming all the columns`
`crime.columns = ['RecNo', 'ReportDate', 'DateOCC', 'TimeOCC', 'Area', 'Area`

In [425]: `# Group by YearMonth and count the number of crimes`
`crime['DateOCC'] = pd.to_datetime(crime['DateOCC'])`
`crime['YearMonth'] = crime['DateOCC'].dt.to_period('M')`
`crime['YearMonth'] = crime['YearMonth'].astype(str)`
`crime`
`crime_trends = crime.groupby('YearMonth').size().reset_index(name='Cr`
`crime_trends`

`# Plotting`
`plt.figure(figsize=(14, 6))`
`sns.lineplot(x='YearMonth', y='Crime Count', data=crime_trends)`
`plt.title('Crime Trends Over Time')`
`plt.xlabel('Year-Month')`
`plt.ylabel('Number of Crimes')`
`plt.xticks(rotation=45)`
`plt.tight_layout()`
`plt.show()`



```
In [426]: # Checking for null entries in our predictors  
print(crime.isnull().sum())
```

```
RecNo          0  
ReportDate     0  
DateOCC        0  
TimeOCC        0  
Area           0  
AreaName       0  
DistrictNo     0  
Part           0  
CrimeCode      0  
CrmDesc        0  
Mocodes        122372  
VictAge        0  
VictSex        116363  
VictRace       116371  
PremiseCd      10  
PremiseDesc    537  
WeaponCd       575065  
WeaponDesc     575065  
Status         0  
StatusDesc     0  
CrimeCd1       11  
CrimeCd2       817200  
CrimeCd3       879589  
CrimeCd4       881713  
Location       0  
CrossStreet    742349  
Lat            0  
Lon            0  
YearMonth      0  
dtype: int64
```

```
In [427]: # Removing records with null values and illogical values
crime = crime.dropna(subset=['Mocodes', 'VictAge', 'VictSex', 'VictRace',
crime['WeaponCd'].fillna(0, inplace=True)


print(crime.isnull().sum())
print(crime.shape)
```

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/pandas/core/series.py:4535: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
downcast=downcast,

RecNo	0
ReportDate	0
DateOCC	0
TimeOCC	0
Area	0
AreaName	0
DistrictNo	0
Part	0
CrimeCode	0
CrmDesc	0
Mocodes	0
VictAge	0
VictSex	0
VictRace	0
PremiseCd	0
PremiseDesc	0
WeaponCd	0
WeaponDesc	452773
Status	0
StatusDesc	0
CrimeCd1	10
CrimeCd2	694377
CrimeCd3	756376
CrimeCd4	758496
Location	0
CrossStreet	641453
Lat	0
Lon	0
YearMonth	0
dtype: int64	
(758558, 29)	

```
In [428]: plt.figure(figsize=(12, 8))
sns.heatmap(crime.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```



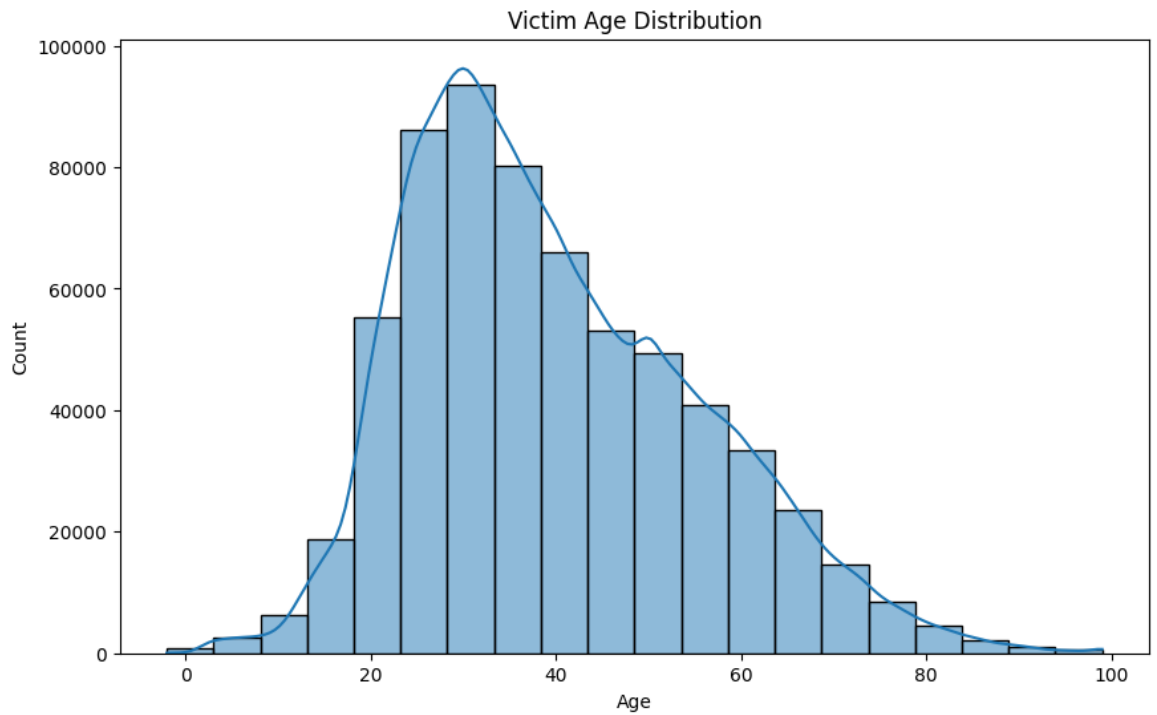
```
In [429]: # Removing unknown records in VictSex and VictRace, Removing 0 in Vict
crime = crime[(crime['VictSex'] != "X") & (crime['VictRace'] != "X")] &
print(crime.shape)
```

```
crime = crime.drop(columns=['CrimeCd2', 'CrimeCd3', 'CrimeCd4', 'Lat',
```

```
(641572, 29)
```




```
In [430]: # Example: Victim Age Distribution
plt.figure(figsize=(10, 6))
sns.histplot(crime["VictAge"], bins=20, kde=True)
plt.title("Victim Age Distribution")
plt.xlabel("Age")
plt.ylabel("Count")
plt.show()
```



```
In [386]: sns.countplot(x='VictSex', data=crime)
plt.title("Victim Gender Distribution")
```

```
Out[386]: Text(0.5, 1.0, 'Victim Gender Distribution')
```



```
In [431]: # Transforming columns - CrimeCode
crime['Severity'] = np.where(crime['CrimeCode'] < 300, 'Severe', 'Non-
crime['Severity'] = pd.Categorical(crime['Severity'])

crime = crime.drop(columns=['CrimeCode', 'CrimeCd1'])
```

```
In [432]: # Transforming columns - VictSex and Weapon
crime['Female'] = np.where(crime['VictSex'] == 'F', 'Yes', 'No')
crime['Female'] = pd.Categorical(crime['Female'])

crime['Weapon'] = np.where(crime['WeaponCd'] == 0, 'No', 'Yes')
crime['Weapon'] = pd.Categorical(crime['Weapon'])

crime = crime.drop(columns=['VictSex', 'WeaponCd'])
```

In [433]:

crime.head()

Out [433]:

	TimeOCC	Area	VictAge	VictRace	PremiseCd	PremiseDesc	YearMonth	Severity	Female
0	2230	3	36	B	501.0	SINGLE FAMILY DWELLING	2020-01	Non- Severe	Yes
1	330	1	25	H	102.0	SIDEWALK	2020-01	Non- Severe	Nc
3	1730	15	76	W	502.0	MULTI-UNIT DWELLING (APARTMENT, DUPLEX, ETC)	2020-01	Non- Severe	Yes
5	30	1	25	H	735.0	NIGHT CLUB (OPEN EVENINGS ONLY)	2020-01	Severe	Yes
6	1315	1	23	H	404.0	DEPARTMENT STORE	2020-01	Non- Severe	Nc

```
In [434]: # Exploratory Data Analysis
severe_exploratory = crime[crime['Severity'] == 'Severe']
severe_exploratory['TimeOCC'] = pd.to_numeric(severe_exploratory['TimeOCC'], errors='coerce')

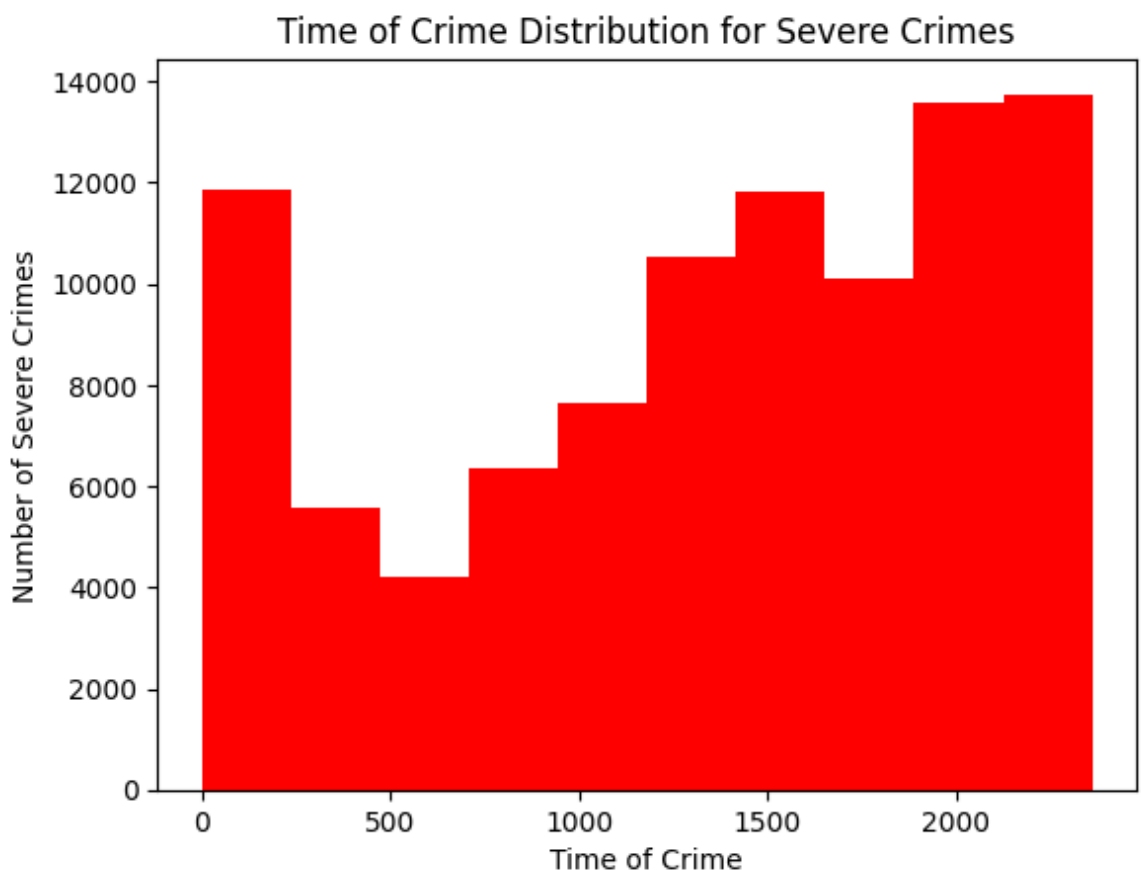
# Time
plt.hist(severe_exploratory['TimeOCC'], color='red')
plt.xlabel('Time of Crime')
plt.ylabel('Number of Severe Crimes')
plt.title('Time of Crime Distribution for Severe Crimes')
plt.show()

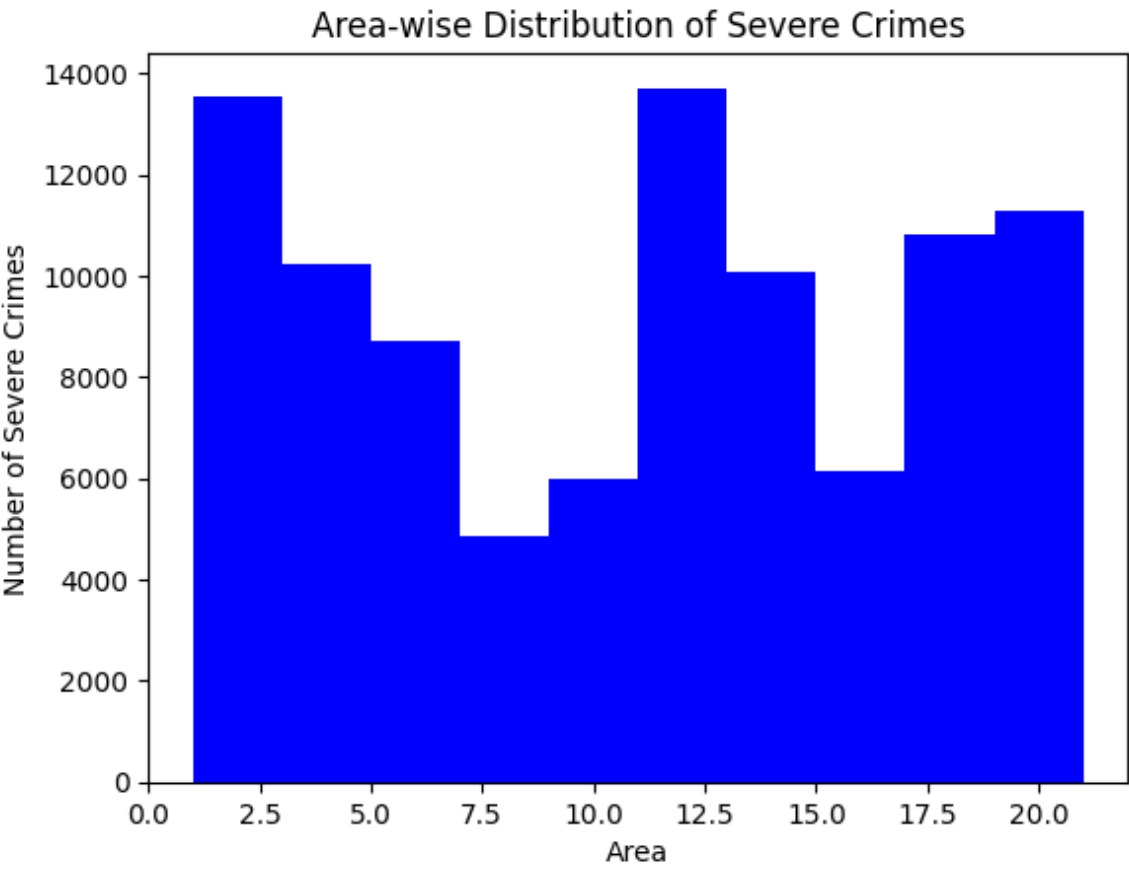
# Area
plt.hist(severe_exploratory['Area'], color='blue')
plt.xlabel('Area')
plt.ylabel('Number of Severe Crimes')
plt.title('Area-wise Distribution of Severe Crimes')
plt.show()
```

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

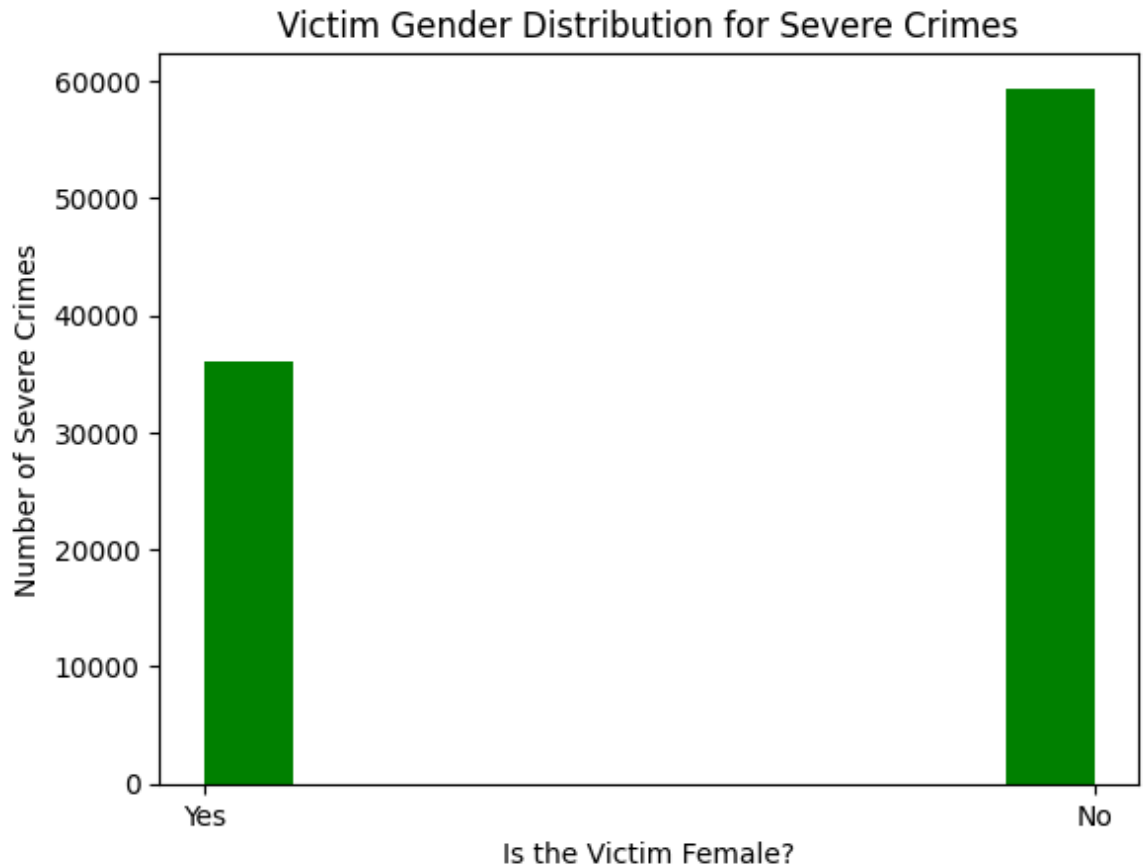
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

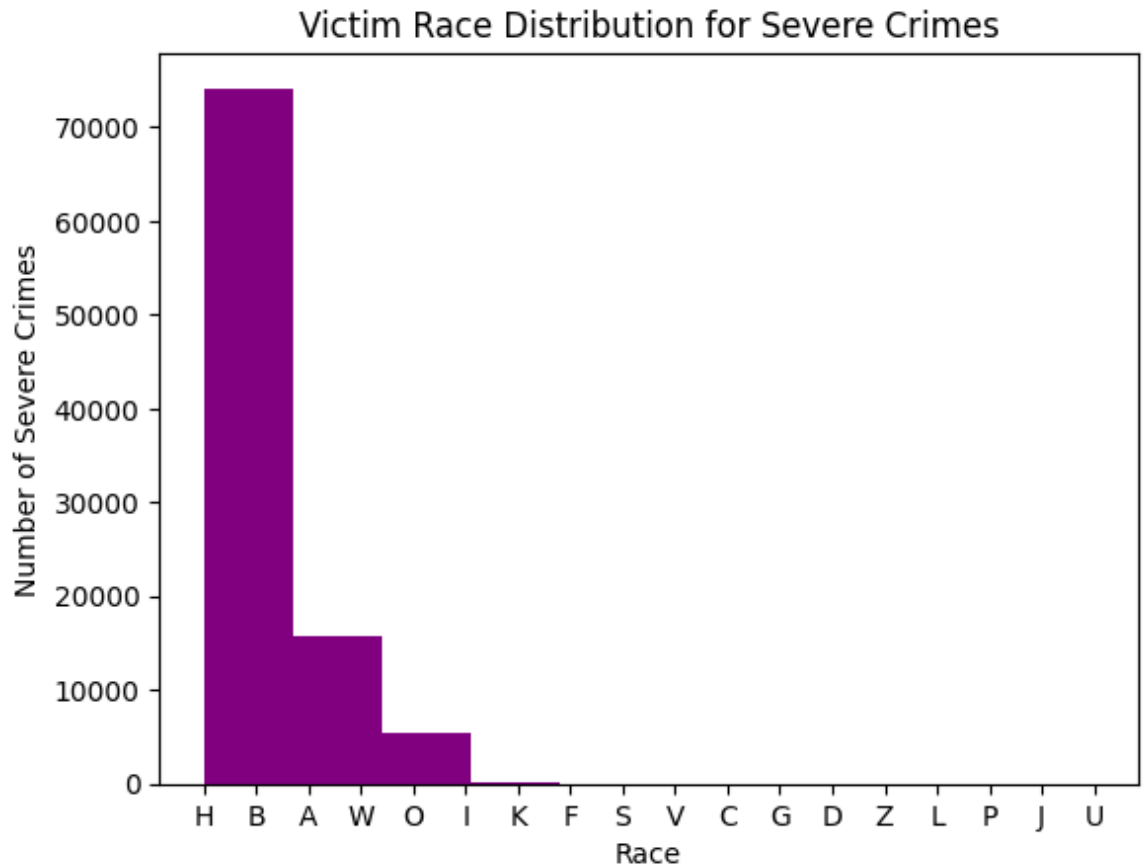




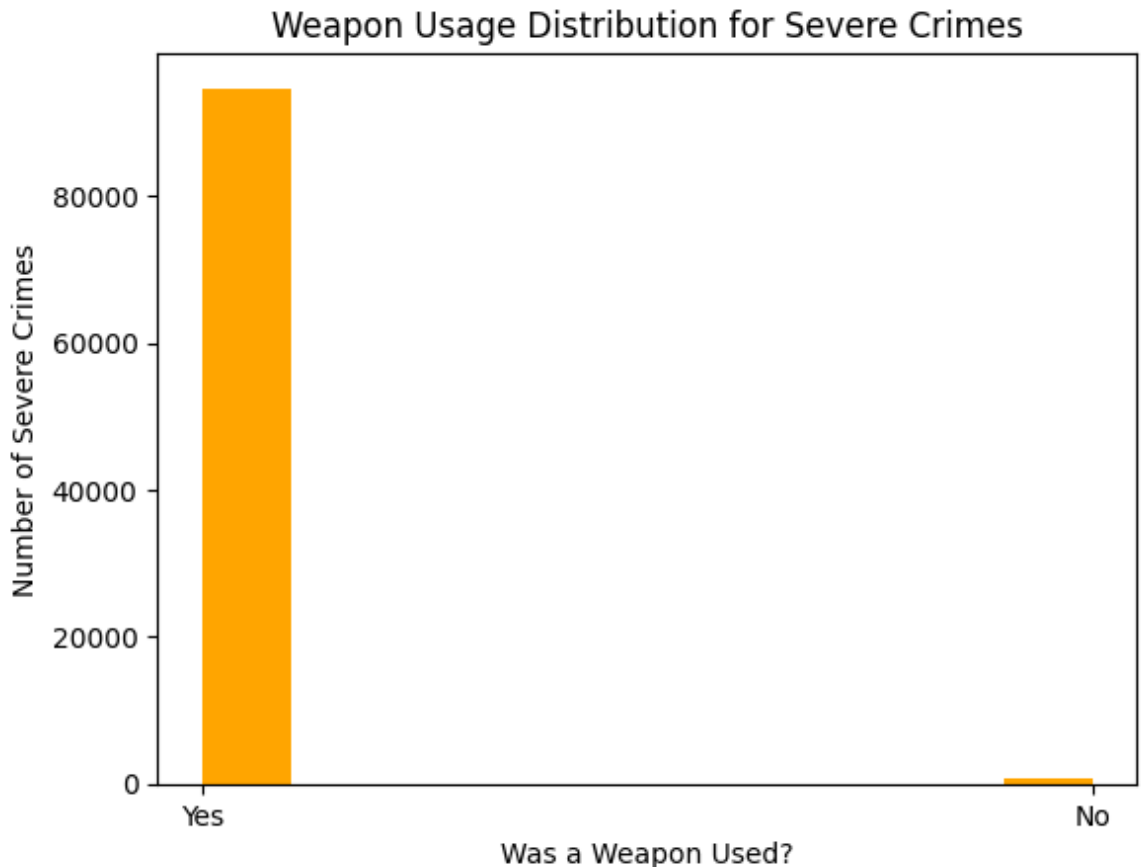
```
In [435]: # Sex
plt.hist(severe_exploratory['Female'], color=['green'])
plt.xlabel('Is the Victim Female?')
plt.ylabel('Number of Severe Crimes')
plt.title('Victim Gender Distribution for Severe Crimes')
plt.show()
```



```
In [436]: # Race
plt.hist(severe_exploratory['VictRace'], color='purple')
plt.xlabel('Race')
plt.ylabel('Number of Severe Crimes')
plt.title('Victim Race Distribution for Severe Crimes')
plt.show()
```



```
In [437]: # Weapon
plt.hist(severe_exploratory['Weapon'], color=['orange'])
plt.xlabel('Was a Weapon Used?')
plt.ylabel('Number of Severe Crimes')
plt.title('Weapon Usage Distribution for Severe Crimes')
plt.show()
```



```
In [446]: # Modelling
X = crime.drop(columns=['Severity', 'YearMonth'])
y = crime['Severity']

categorical_features = ['Area', 'VictRace', 'PremiseDesc', 'Female', '
numeric_features = ['TimeOCC', 'VictAge', 'PremiseCd']
```

```
In [447]: # Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.

# Preprocessing: One-hot encoding categorical features and scaling num
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
        ('cat', OneHotEncoder(), categorical_features)
    ],
    remainder='passthrough'
)
```



```
In [449]: X_train_preprocessed = preprocessor.fit_transform(X_train)
X_test_preprocessed = preprocessor.transform(X_test)

# Decision Tree model
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train_preprocessed, y_train)
dt_pred = dt_model.predict(X_test_preprocessed)
```

```
In [450]: # Printing the classification report
print("Decision Tree Classifier Report:")
print(classification_report(y_test, dt_pred))
```

```
Decision Tree Classifier Report:
              precision    recall  f1-score   support

   Non-Severe      0.90      0.90      0.90     109271
     Severe       0.44      0.44      0.44      19044


   accuracy              0.83     128315
  macro avg       0.67      0.67      0.67     128315
weighted avg       0.83      0.83      0.83     128315
```

```
In [451]: # Predict probabilities for positive class
y_prob = dt_model.predict_proba(X_test_preprocessed)[: , 1]

# Calculate false positive rate (FPR), true positive rate (TPR), and t
fpr, tpr, thresholds = roc_curve(y_test, y_prob, pos_label='Severe')

# Calculate Area Under the Curve (AUC)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)'
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



```
In [452]: # Logistic Regression
from sklearn.linear_model import LogisticRegression
lr_model = LogisticRegression(max_iter=1000)
lr_model.fit(X_train_preprocessed, y_train)
lr_pred = lr_model.predict(X_test_preprocessed)
```

```
In [453]: print("Logistic Regression Classifier Report:")
print(classification_report(y_test, lr_pred))
```

Logistic Regression Classifier Report:

	precision	recall	f1-score	support
Non-Severe	0.90	0.96	0.93	109271
Severe	0.61	0.38	0.46	19044
accuracy			0.87	128315
macro avg	0.75	0.67	0.70	128315
weighted avg	0.86	0.87	0.86	128315

```
In [454]: # Import roc_auc_score
from sklearn.metrics import roc_auc_score


# Convert categorical labels to binary labels
y_test_binary = y_test.map({'Non-Severe': 0, 'Severe': 1})

# Calculate ROC curve
fpr, tpr, thresholds = roc_curve(y_test_binary, lr_probs)

# Calculate AUC
lr_auc = roc_auc_score(y_test_binary, lr_probs)

# Plot ROC curve
plt.plot(fpr, tpr, marker='.')
plt.title('ROC Curve for Logistic Regression')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()

print("AUC for Logistic Regression:", lr_auc)
```



AUC for Logistic Regression: 0.8967405408436027

In []: